

DDPM sur MNIST

Arthur Hamard, Arthur Guelennoc, Léo Sambrook, Etienne Senigout,
Pierre Braud, Paul Pazart

Introduction

Le Projet en un Clin d'œil

- **Maîtriser le cycle :** Destruction \leftrightarrow Reconstruction
- **Technique :** Denoising Diffusion Probabilistic Models (DDPM)
- **Dataset :** MNIST (60k chiffres)



Partie 1 : Théorie - Le Principe

Forward vs Reverse

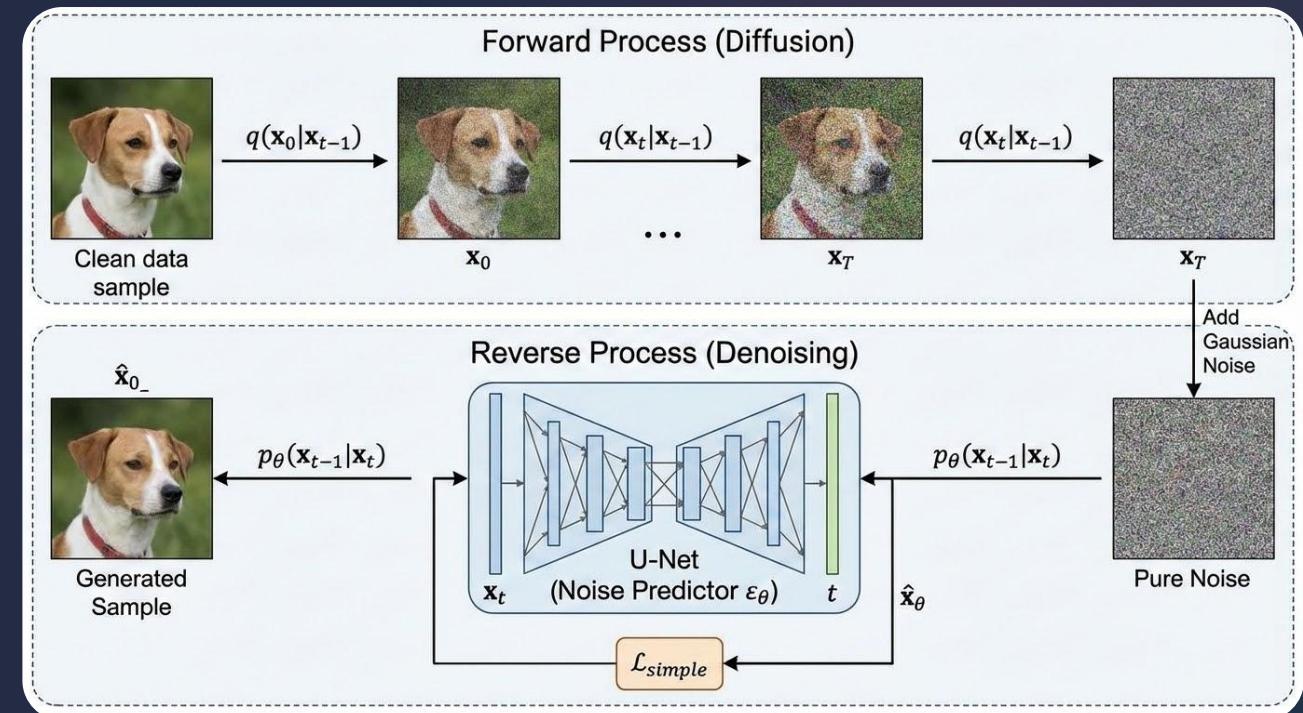
L'idée centrale est de détruire l'information progressivement pour ensuite apprendre à la reconstruire.

- **Processus Forward (q):**

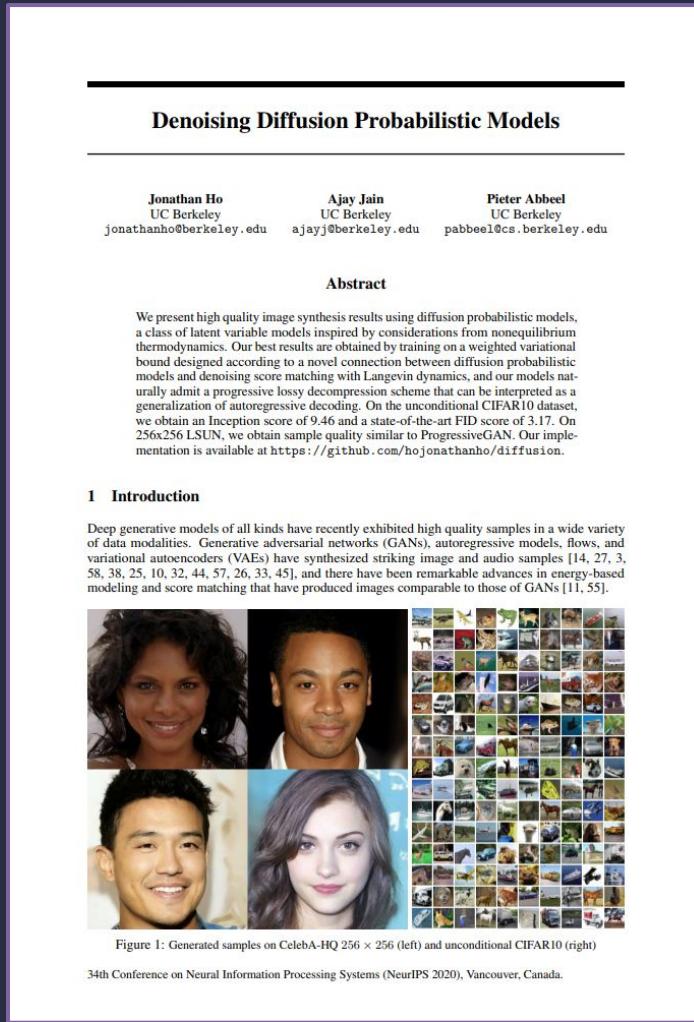
Une chaîne de Markov fixe qui ajoute du bruit gaussien. À T , l'image est un bruit pur $N(0, 1)$.

- **Processus Reverse (p_θ):**

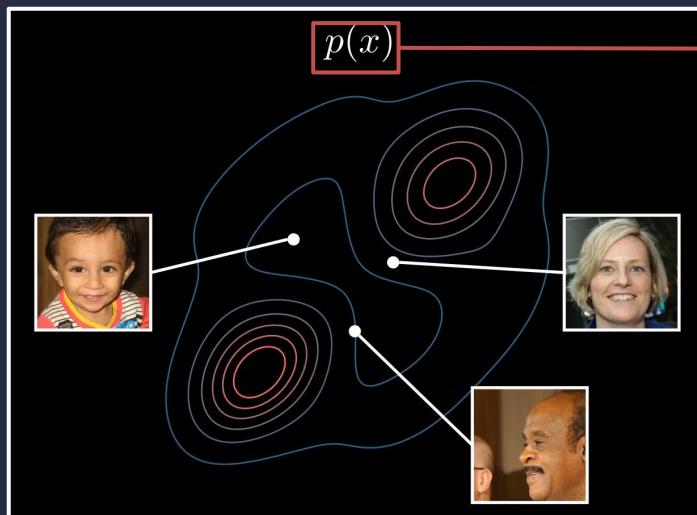
Un réseau de neurones apprend à inverser ce processus, étape par étape.



Formulation Mathématique : Intro (accrochez vous!)

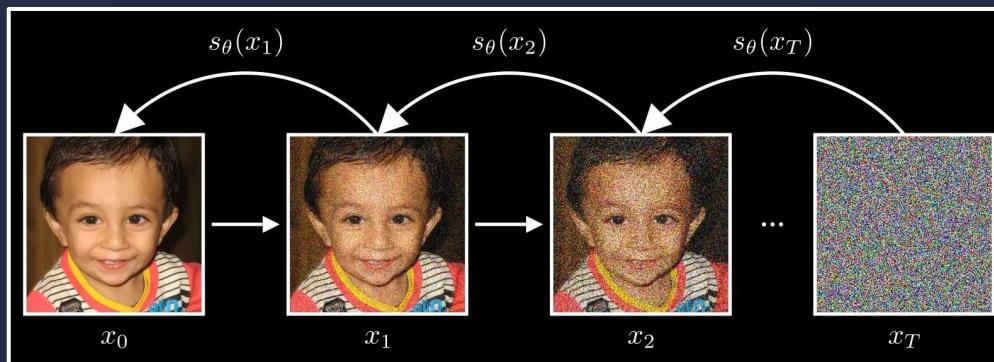


- a “introduit” la diffusion
- 20K+ citations



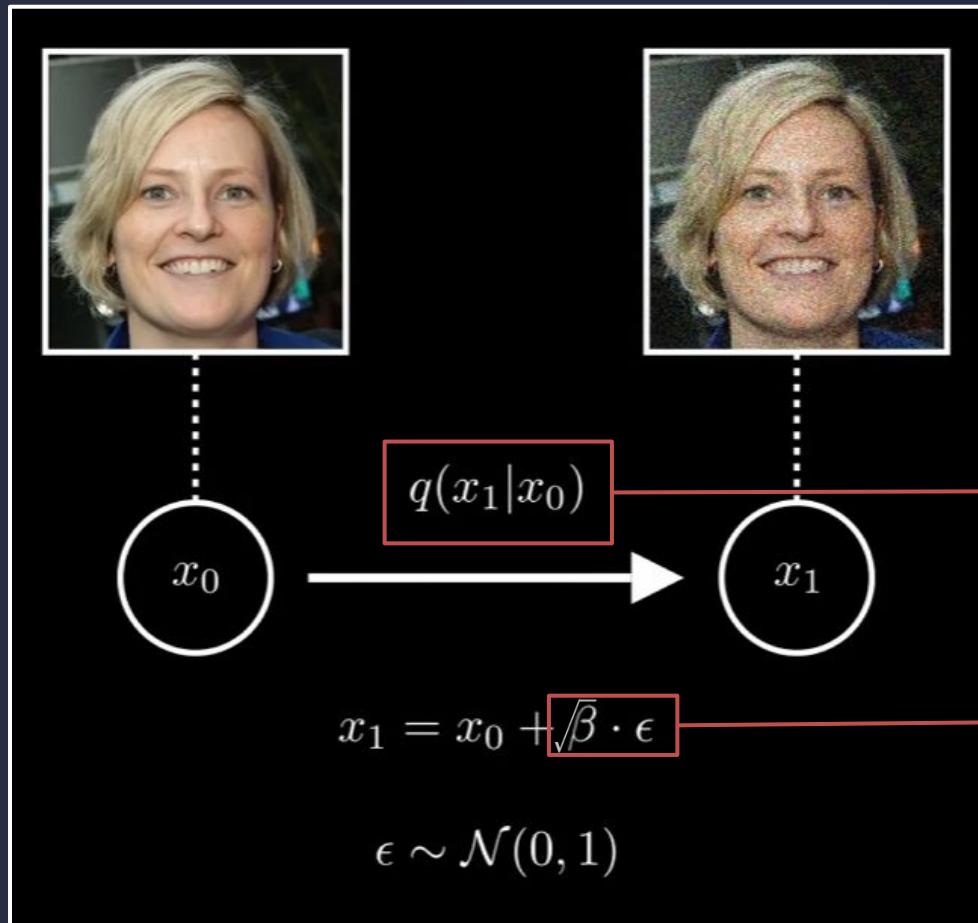
distribution des données

- dataset d’images
- p inconnu
- $p(x) \sim$ plausibilité de l’image
- image du dataset $\Rightarrow p(x)$ élevé
- image absurde $\Rightarrow p(x)$ petit
- générer images similaires



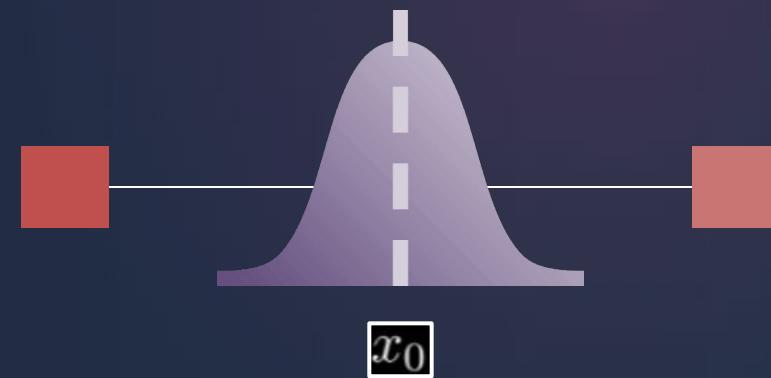
- modélisation impossible
- corruption graduelle
- bruit pur
- x_0, x_1, \dots nos images
- reverse
- forward = +, backward = -

Formulation Mathématique : Forward



distribution conditionnelle

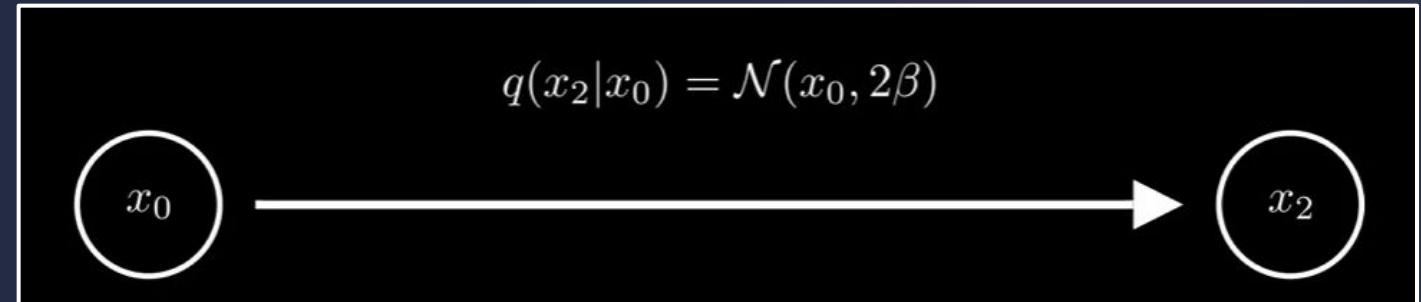
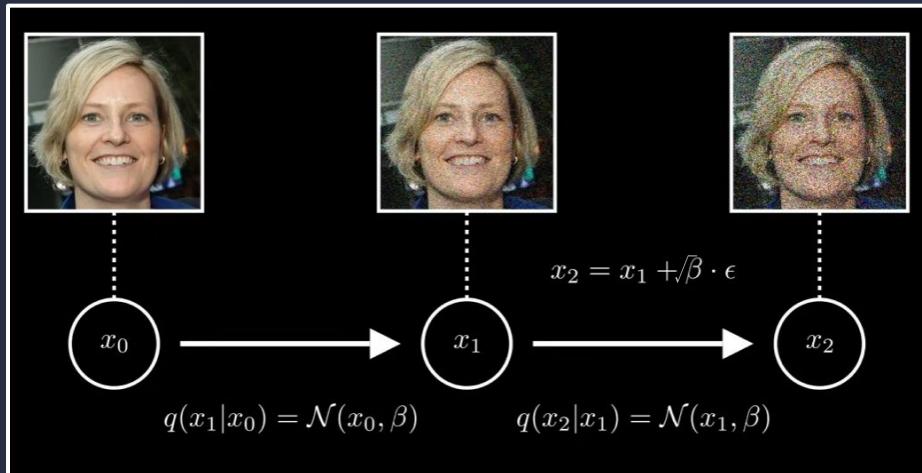
bruit
Gaussien
contrôlé par
un scalaire β



$$q(x_1|x_0) = \mathcal{N}(x_0, \beta)$$

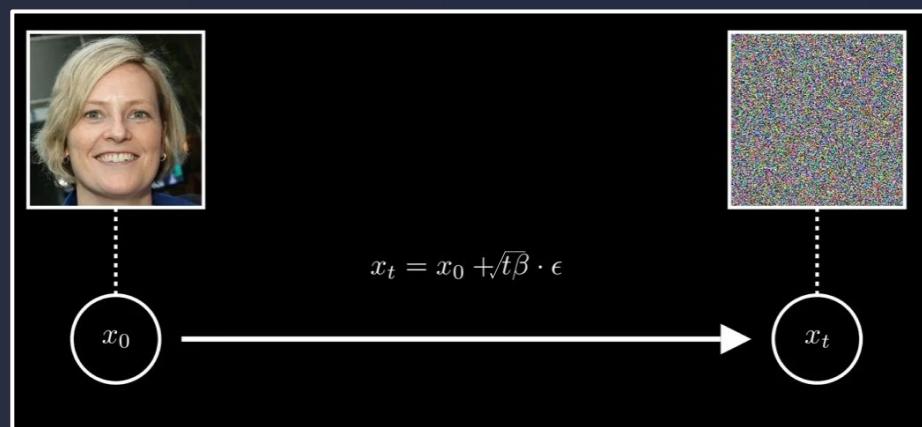
distrib de x_1 = Gaussienne centrée sur x_0 ,
bruit de variance β

Formulation Mathématique : Forward



- pareil que sampler la Gaussienne de variance 2β car $x_2 = x_0 + \sqrt{\beta} \epsilon_1 + \sqrt{\beta} \epsilon_2$ $x_2 = x_0 + \sqrt{\beta} (\epsilon_1 + \epsilon_2)$ or $\epsilon_1 + \epsilon_2 \sim \mathcal{N}(0, 2)$
- $x_2 = x_0 + \sqrt{2\beta} \epsilon$

- on répète l'opération



- $\sqrt{t\beta}$ de bruit
- aka $q(x_t | x_0) = \mathcal{N}(x_0, t\beta)$ à t ,
- $q(x_t | x_0) =$
distribution Gaussienne centrée sur x_0 , de variance $t\beta$

Formulation Mathématique : Forward (2)

$$q(x_t | x_0) = \mathcal{N}(x_0, t\beta)$$

Explosion de la variance de diffusion



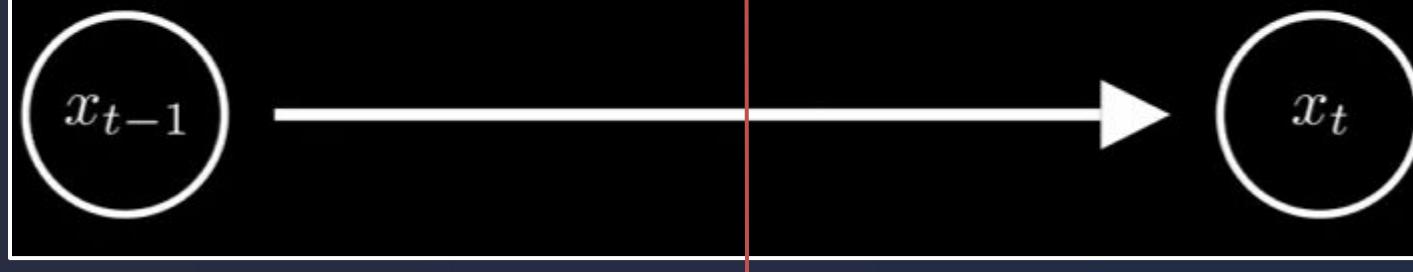
$$q(x_t|x_0) \xrightarrow[t \rightarrow \infty]{} \mathcal{N}(0, 1)$$

Simplification des calculs

- Forward
- Backward
- Base de génération

Formulation Mathématique : Forward (2)

$$q(x_t|x_{t-1}) = \sqrt{1 - \beta} x_{t-1} + \sqrt{\beta} \cdot \epsilon$$

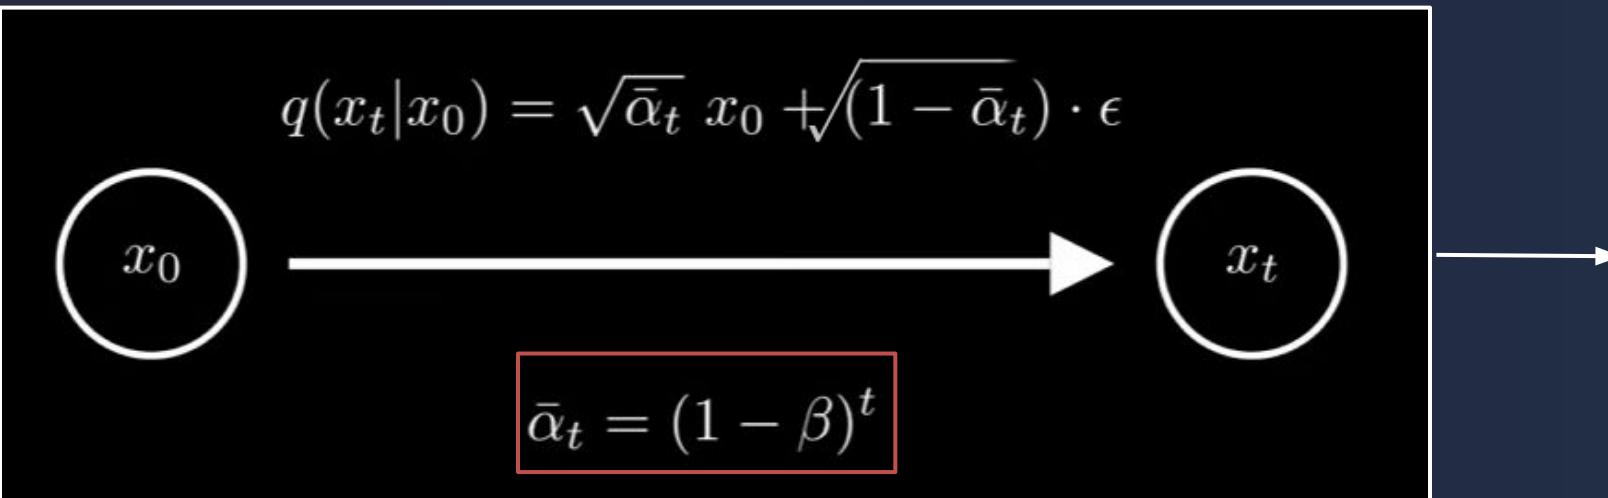


coefficient pour diminuer x

pourquoi ?

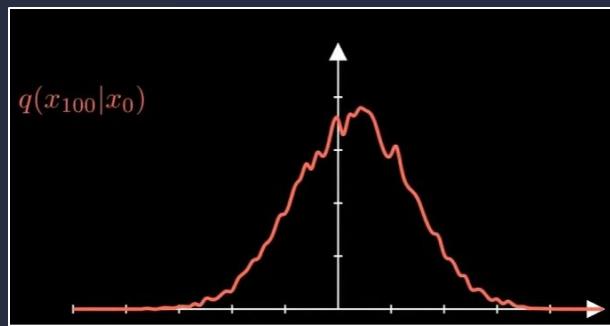
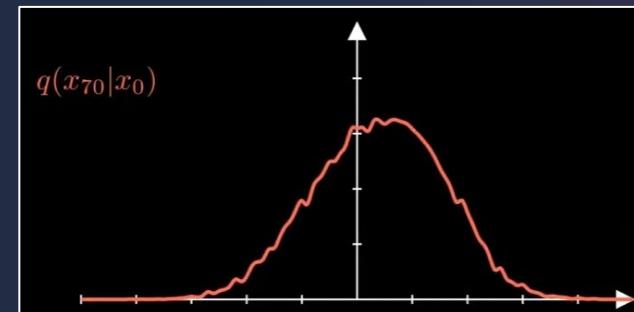
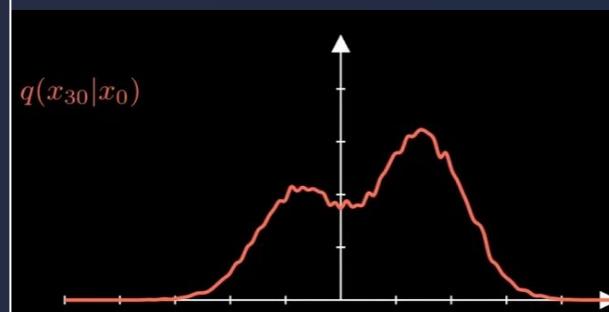
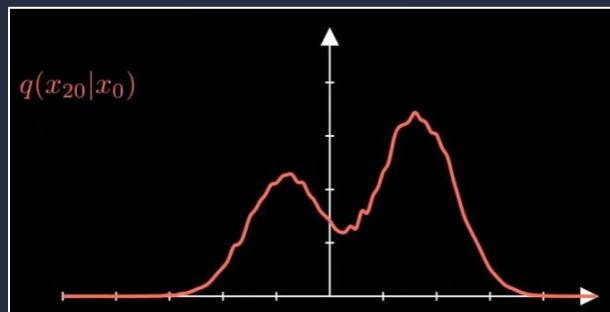
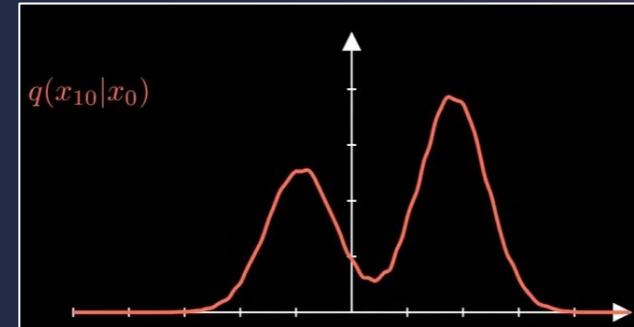
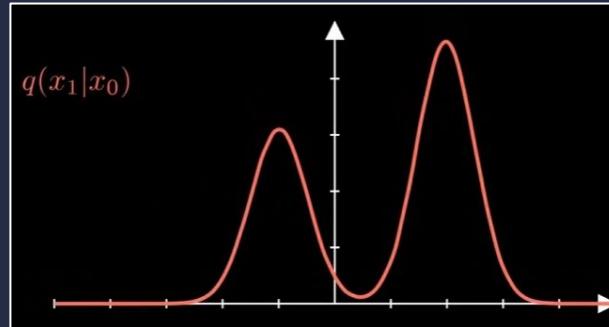
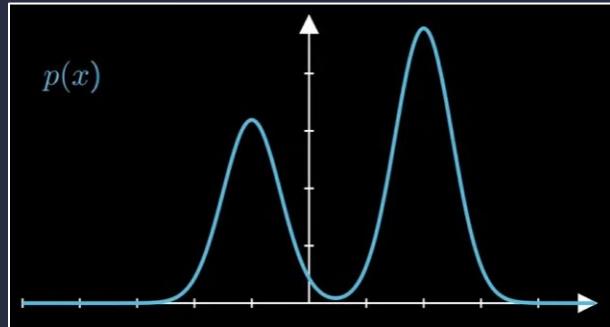
- auteurs de DDPM
- non-expliqué
- simple

$$q(x_t|x_0) = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{(1 - \bar{\alpha}_t)} \cdot \epsilon$$



moyenne $\rightarrow 0$
variance $\rightarrow 1$
temps $\rightarrow \infty$

Formulation Mathématique : Forward (2)



distribution des données —————→ **distribution centrée réduite**

pas besoin d'itérer, calcul de t=200 instantané

Formulation Mathématique : Forward (2)

$$q(x_t|x_0) = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \cdot \epsilon$$



$$\bar{\alpha}_t = (1 - \beta)^t$$

- **β n'est pas fixe**

- **β scheduler**

- **$0 < \beta < 1$**

$$q(x_t|x_0) = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \cdot \epsilon$$

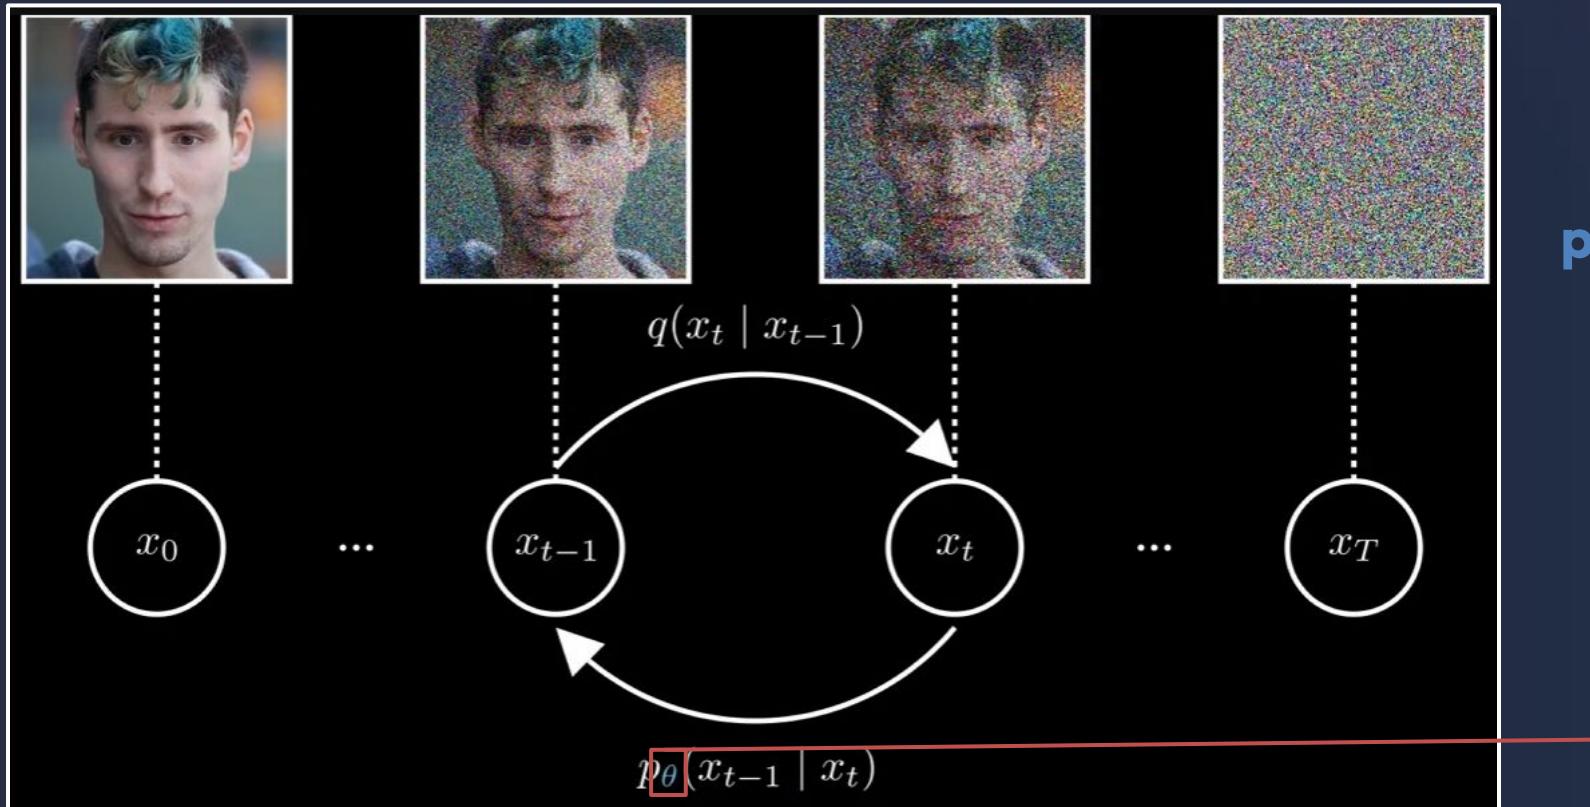


$$\bar{\alpha}_t = \prod_{i=1}^t (1 - \beta_i)$$

expression du papier

Formulation Mathématique : Backward

Comment inverser le procédé ?

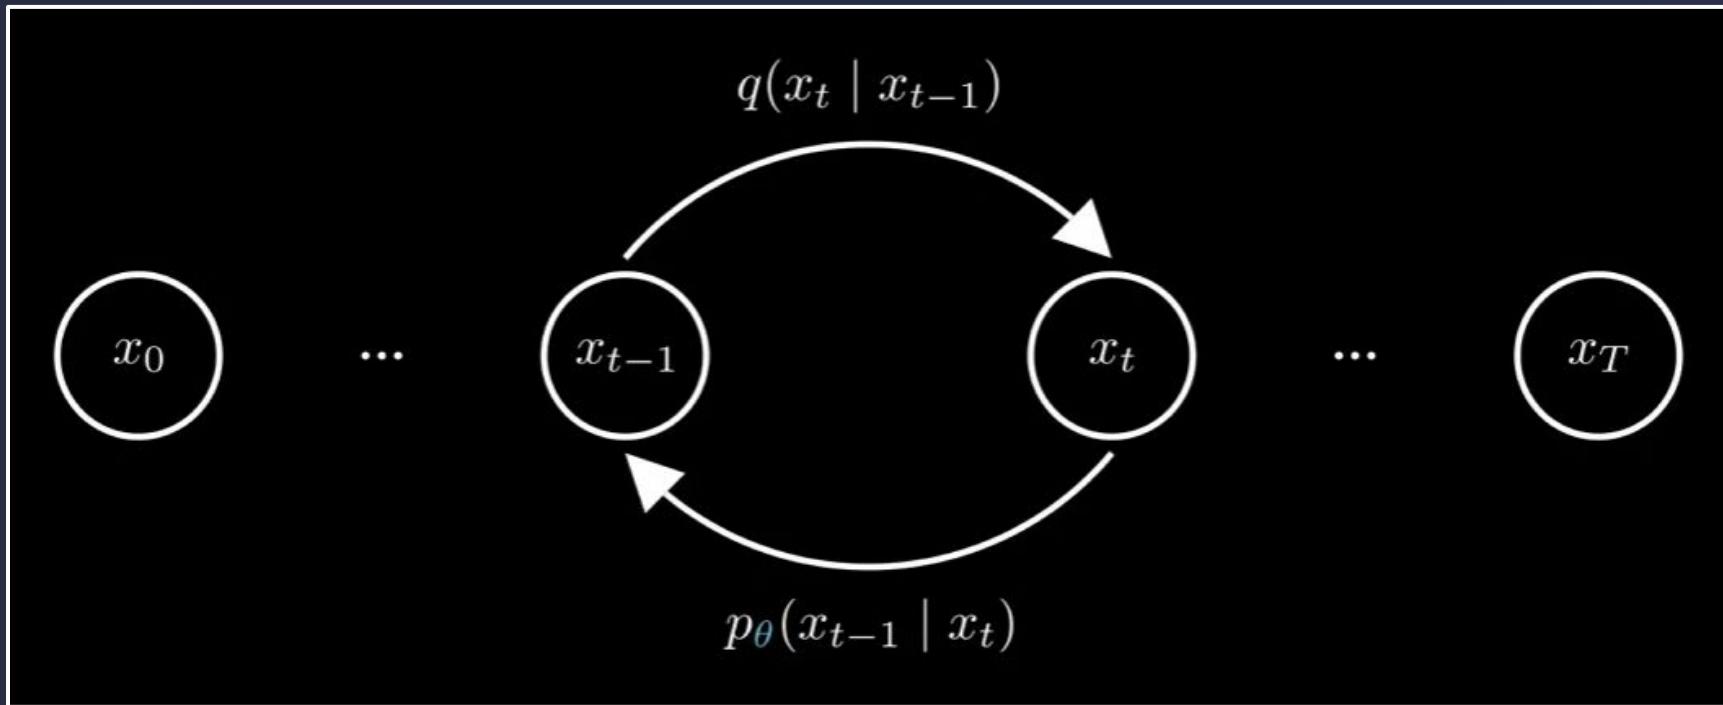


paramètres fixe en forward

mais en backward

on veut le meilleur θ
= poids du réseau

Formulation Mathématique : Backward (entraînement)



approche bayésienne

on veut minimiser

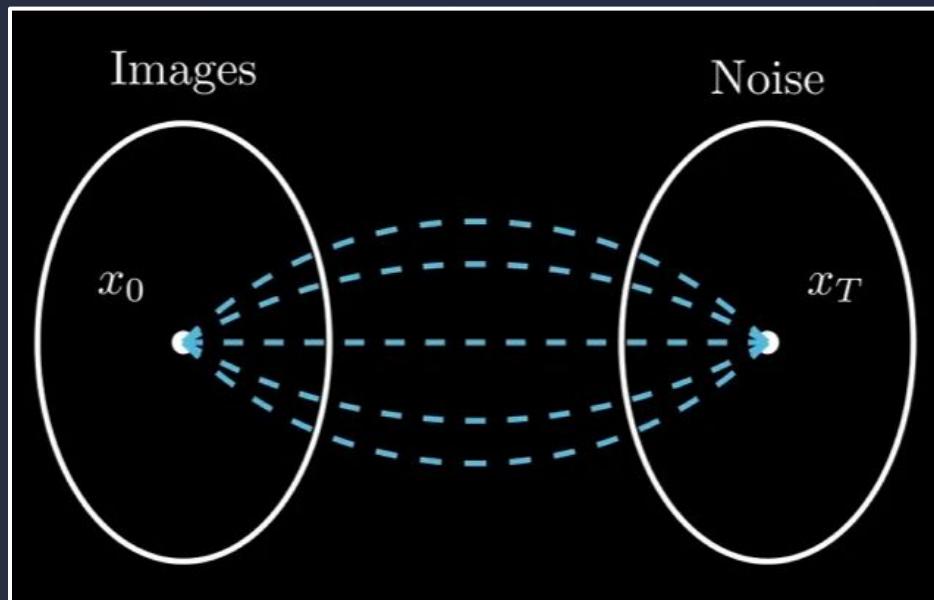
$$-\log p_\theta(x_0)$$

negative log-likelihood

Formulation Mathématique : Loss

$$-\log p_{\theta}(x_0) = -\log \int p_{\theta}(x_{0:T}) dx_{1:T}$$

intégrale de la probabilité conjointe sur toutes les variables



“chemins” qui connectent image à leur bruit respectif
negative log likelihood nécessite addition de tous ces chemins possibles

beaucoup trop de chemins possibles

Formulation Mathématique : Loss (2)

$$-\log p_{\theta}(x_0) = -\log \int p_{\theta}(x_{0:T}) dx_{1:T}$$

↓
multiplication + division par la densité
du procédé forward

$$-\log p_{\theta}(x_0) = -\log \int q(x_{1:T} \mid x_0) \frac{p_{\theta}(x_{0:T})}{q(x_{1:T} \mid x_0)} dx_{1:T}$$

$$-\log p_{\theta}(x_0) = -\log \mathbb{E}_{q(x_{1:T} \mid x_0)} \left[\frac{p_{\theta}(x_{0:T})}{q(x_{1:T} \mid x_0)} \right]$$

↓
inégalité de Jensen
negative log convexe

$$-\log p_{\theta}(x_0) \leq -\mathbb{E}_{q(x_{1:T} \mid x_0)} \left[\log \frac{p_{\theta}(x_{0:T})}{q(x_{1:T} \mid x_0)} \right]$$

ELBO
Evidence Lower Bound

$$-\mathbb{E}_{q(x_{1:T} \mid x_0)} \left[\log \frac{p_{\theta}(x_{0:T})}{q(x_{1:T} \mid x_0)} \right]$$

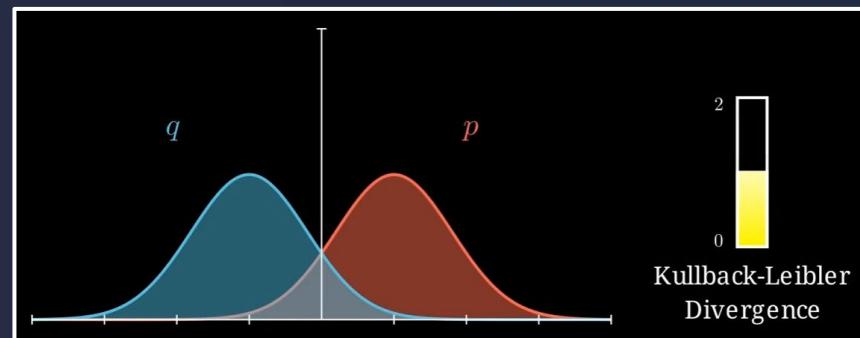
Formulation Mathématique : Loss (3)

$$-\mathbb{E}_{q(x_{1:T} | x_0)} \left[\log \frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)} \right]$$

calculs expliqués en annexe du papier

$$-\mathbb{E}_q \left[D_{\text{KL}}(q(x_T | x_0) || p(x_T)) + \sum_{t>1} D_{\text{KL}}(q(x_{t-1} | x_t, x_0) || p_\theta(x_{t-1} | x_t)) - \log p_\theta(x_0 | x_1) \right]$$

Kullback-Leibler divergence



distance entre
distributions de probabilité

quand 0 ~ pareil

Formulation Mathématique : Loss (4)

$$-\mathbb{E}_q \left[D_{\text{KL}}(q(x_T | x_0) || p(x_T)) + \sum_{t>1} D_{\text{KL}}(q(x_{t-1} | x_t, x_0) || p_{\theta}(x_{t-1} | x_t)) - \log p_{\theta}(x_0 | x_1) \right]$$

ne dépend pas des paramètres

Négligeable

Formulation Mathématique : Loss (5)

$$-\mathbb{E}_q \left[\sum_{t>1} D_{\text{KL}}(q(x_{t-1} | x_t, x_0) || p_\theta(x_{t-1} | x_t)) \right]$$

vrai posterior

posterior approximé

$$q(x_{t-1} | x_t, x_0) \propto q(x_t | x_{t-1}) q(x_{t-1} | x_0) = \mathcal{N}(x_{t-1}; \tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$$
$$p_\theta(x_{t-1} | x_t) = \mathcal{N}(\mu_\theta, \sigma_t)$$
$$\tilde{\beta}_t = \sigma_t$$

$$D_{\text{KL}}(q(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) = D_{\text{KL}}(\mathcal{N}(x_{t-1}; \tilde{\mu}_t, \tilde{\beta}_t I) \| \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \tilde{\beta}_t I))$$
$$= \frac{1}{2\tilde{\beta}_t} \|\tilde{\mu}_t - \mu_\theta\|^2$$

Formulation Mathématique : Loss (6)

$$L = \mathbb{E}_{x_0, \epsilon} \left[\sum_{t=1}^T \frac{1}{2\tilde{\beta}_t} \left\| \tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t) \right\|^2 \right]$$

$$\begin{aligned}\tilde{\mu}_t(x_t, x_0) &= \frac{\sqrt{\alpha_{t-1}}\beta_t}{1-\bar{\alpha}_t}x_0 + \frac{\sqrt{\alpha_t}(1-\bar{\alpha}_{t-1})}{1-\bar{\alpha}_t}x_t \\ \tilde{\beta}_t &= \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t\end{aligned}$$

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1-\bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

$$\begin{aligned}\tilde{\mu}_t(x_t, \epsilon) &= \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon \right) \\ \mu_\theta(x_t, t) &= \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right)\end{aligned}$$

Formulation Mathématique : Recap

$$-\log p_\theta(x_0) = -\log \int p_\theta(x_{0:T}) dx_{1:T}$$

$$-\log p_\theta(x_0) \leq -\mathbb{E}_{q(x_{1:T} | x_0)} \left[\log \frac{p_\theta(x_{0:T})}{q(x_{1:T} | x_0)} \right]$$

$$L = \mathbb{E}_{x_0, \epsilon} \left[D_{KL}(q(x_T | x_0) \| p(x_T)) + \sum_{t=2}^T D_{KL}(q(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) - \log p_\theta(x_0 | x_1) \right]$$

$$L = \mathbb{E}_{x_0, \epsilon} \left[\sum_{t=1}^T D_{\text{KL}}(q(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) \right]$$

$$L = \mathbb{E}_{x_0, \epsilon} \left[\sum_{t=1}^T \frac{1}{2\tilde{\beta}_t} \left\| \tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t) \right\|^2 \right]$$

Formulation Mathématique : Recap

$$L = \mathbb{E}_{x_0, \epsilon} \left[\sum_{t=1}^T \frac{1}{2\tilde{\beta}_t} \left\| \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon \right) - \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) \right\|^2 \right]$$

$$L = \mathbb{E}_{x_0, \epsilon} \left[\sum_{t=1}^T \frac{1}{2\tilde{\beta}_t \alpha_t} \left\| \frac{\beta_t}{\sqrt{1-\bar{\alpha}_t}} (\epsilon - \epsilon_\theta(x_t, t)) \right\|^2 \right]$$

$$L = \mathbb{E}_{x_0, \epsilon} \left[\sum_{t=1}^T \frac{\beta_t^2}{2\tilde{\beta}_t \alpha_t (1-\bar{\alpha}_t)} \left\| \epsilon - \epsilon_\theta(x_t, t) \right\|^2 \right]$$

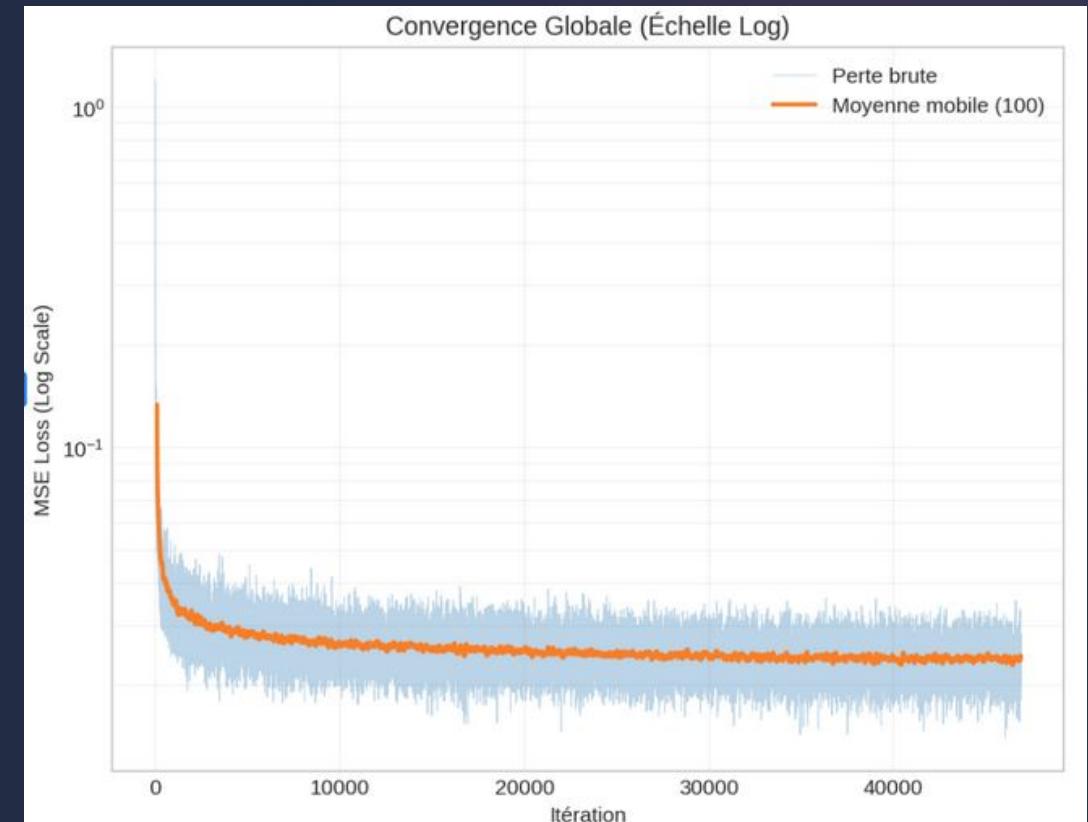
$$L \propto \mathbb{E}_{x_0, \epsilon} \left[\sum_{t=1}^T \left\| \epsilon - \epsilon_\theta(x_t, t) \right\|^2 \right]$$

$$L = \mathbb{E}_{x_0, \epsilon, t} \left[\left\| \epsilon - \epsilon_\theta(x_t, t) \right\|^2 \right], \quad t \sim \text{Uniform}\{1, \dots, T\}$$

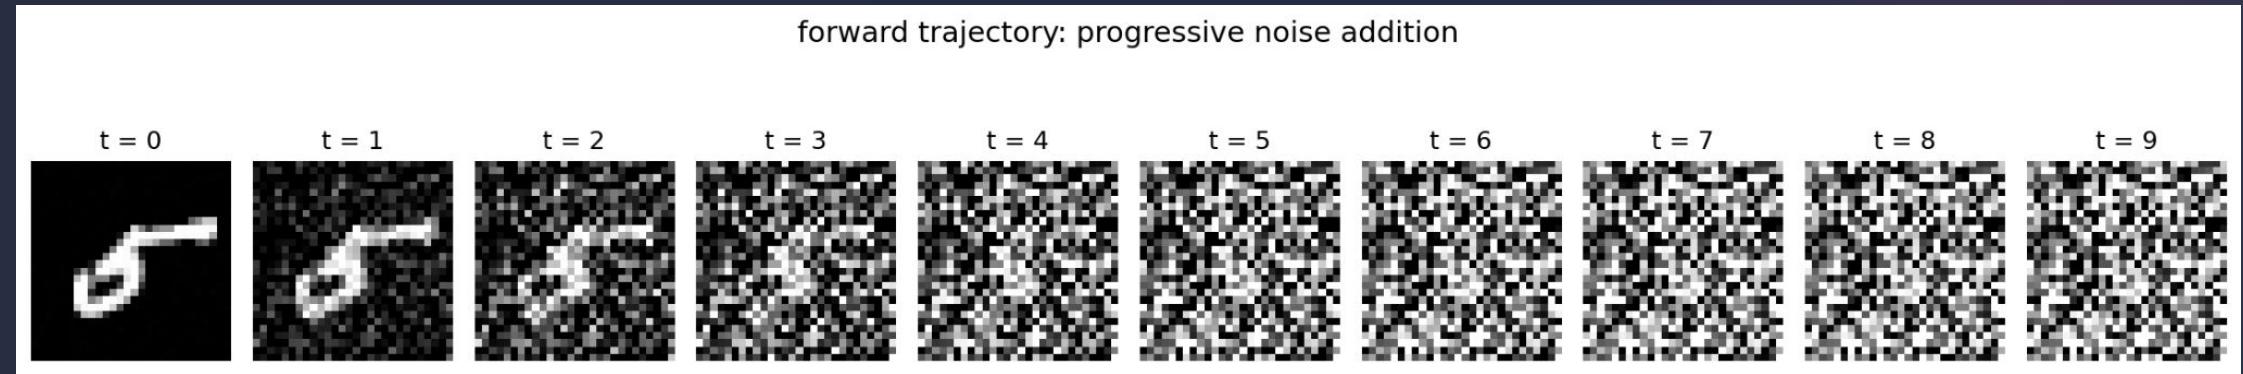
L'Entraînement (Context & Loss)

Convergence et Stabilité

- **Architecture :** U-Net Light (~1.7M params)
- **Dataset :** MNIST (28x28)
- **Steps de Diffusion :** T=1000



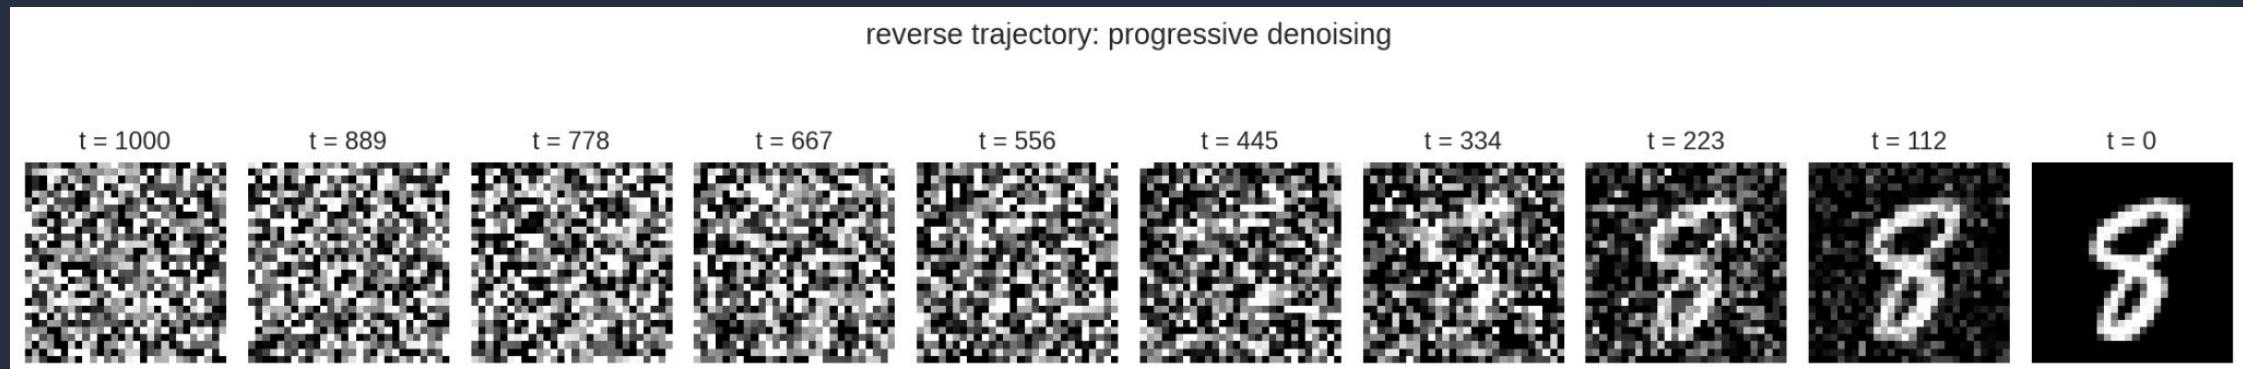
Le Cycle de Diffusion : Destruction & Reconstruction



Forward Process (q) : Ajout de bruit



reverse trajectory: progressive denoising



Reverse Process ($p\theta$) : Génération



Preuve de Qualité (Métriques)

Évaluation réalisée sur **1000 échantillons** générés.

22.58

FID SCORE

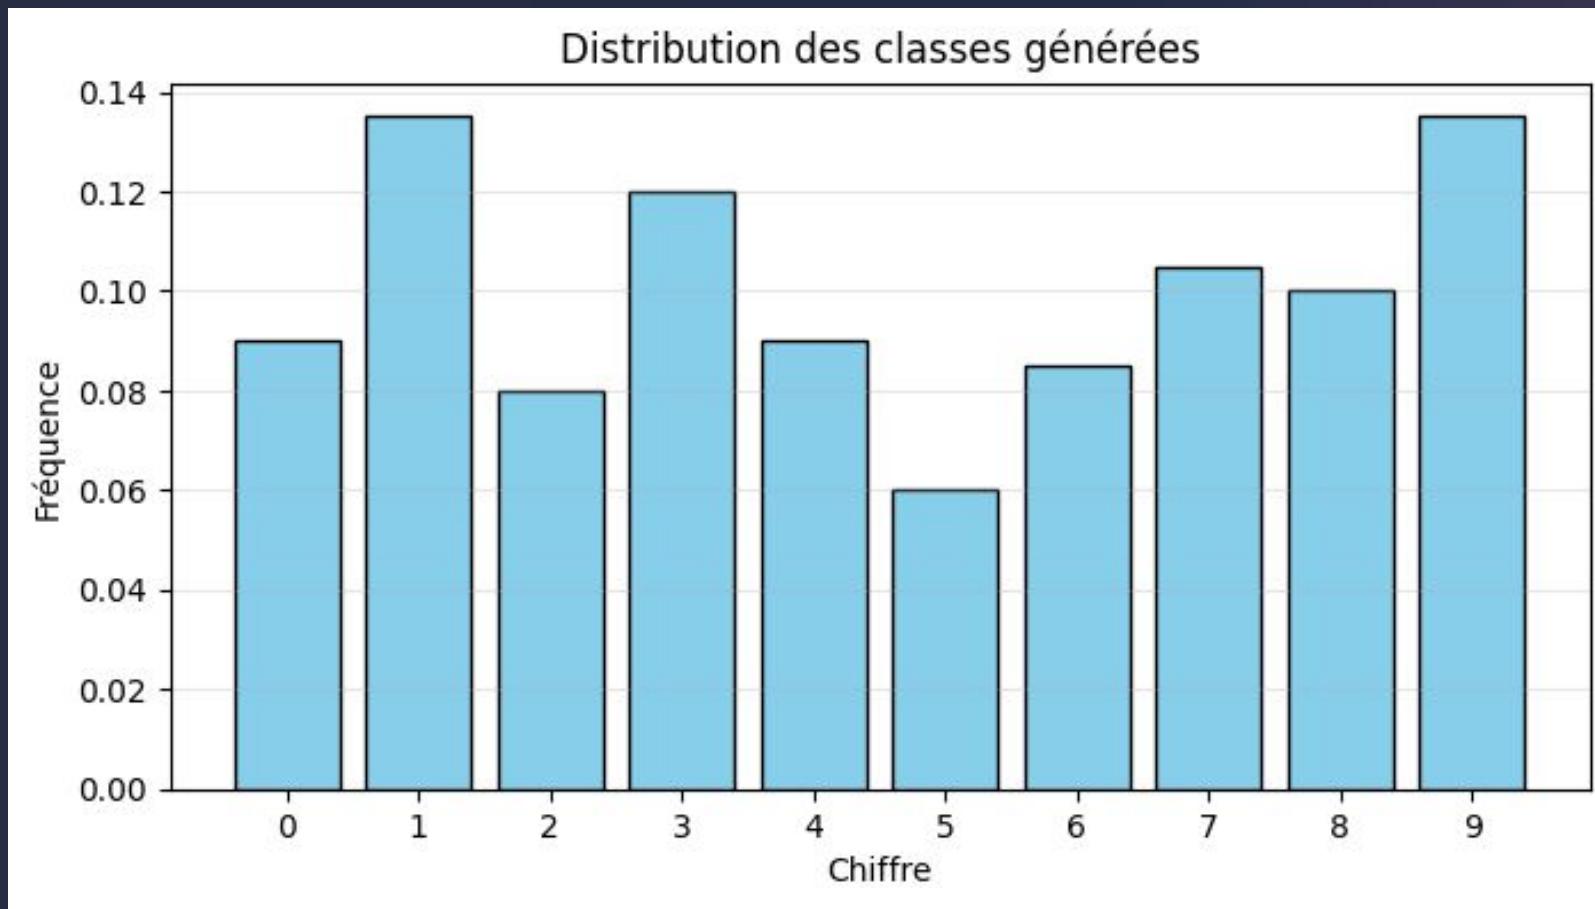
Similarité perceptuelle avec MNIST.

93.05

CONFiance MOYENNE

Les chiffres sont clairs et non ambigus.

Preuve de Qualité (Métriques)



Pas de "Mode Collapse" majeur. Toutes les classes sont générées.

DEMO

Conclusion

- ✓ Objectif atteint : Implémentation fonctionnelle.
- ✗ Performance : Entraînement stable, FID ~22.
- Futur : Explorer DDIM et Latent diffusion.